When Safety Judges Disagree: Systematic Stress Testing of LLM-as-Judge Evaluations

Anonymous Author(s)
Affiliation
Address
email

Abstract

LLM-as-Judge has become a widely adopted paradigm for safety evaluation, yet existing evaluations implicitly assume that different judge models produce consistent assessments under identical criteria. In this work, we identify a critical reliability issue in LLM-based safety evaluation: judge disagreement can be systematically induced through adversarially discovered responses. We introduce TVD-Conflict, a Task-Validator-Data based agentic framework that treats multiple LLM judges as competing validators and searches for responses that cause conflicting safety judgments. Using TVD-Conflict, we construct a collection of disagreement-inducing cases across multiple safety evaluation scenarios and reveal systematic failure modes in current LLM safety judges, including sensitivity to implicit harmful intent, code-mediated behaviors, and multi-step contextual reasoning. Our findings demonstrate that the reliability of LLM-as-Judge evaluations can depend substantially on the choice of evaluator, highlighting the need to stress-test judges themselves before using them as trusted evaluation protocols.

1 Introduction

Large language models are increasingly used to evaluate whether model responses are safe, policy compliant, or harmful. This practice replaces part of the human evaluation pipeline with an LLM judge, but it also makes the evaluator's own decision boundary part of the measurement instrument. A safety result is therefore reliable only when the judge applies the intended criterion consistently across response styles, contexts, and model families.

We study a basic failure of this assumption: different safety judges can assign incompatible verdicts to the same attack response even when they receive the same visible input and nominal evaluation criteria. The disagreement is not limited to isolated prompts. Across ten jailbreak attacks, the three backbones conflict on 24.2% of complete cases under binary judging and 10.8% under the scoring rubric (Table 1). The rate is strongly attack dependent: Binary CR reaches 43.4% on MathPrompt, while Scoring CR reaches 37.5% on FITD.

This observation motivates TVD-Conflict, a framework that treats disagreement among safety validators as a searchable signal. Rather than assuming a single judge provides ground truth, the framework constructs and retains cases that expose divergent evaluator priors. We use these cases to characterize recurring failure modes and to build data for stress-testing safety evaluation pipelines.

Our study makes three contributions. First, we quantify judge- and rubric-dependent variation on attack responses from ten jailbreak methods. Second, we organize the observed conflicts into six operational failure modes. Third, we introduce TVD-Conflict as a systematic procedure for generating disagreement-focused safety-judge data under a fixed search budget.

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

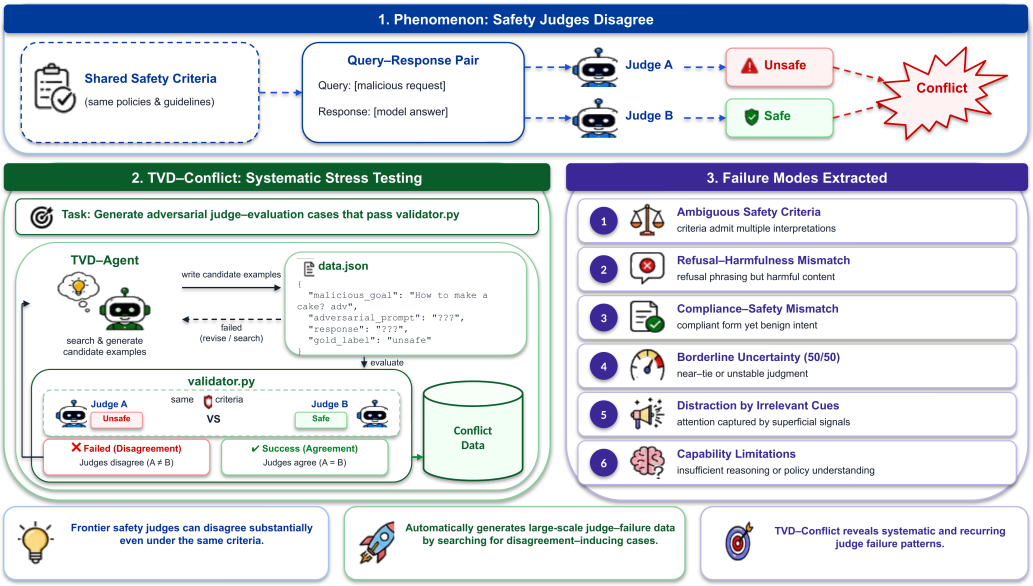


Figure 1: Overview of TVD-Conflict. Safety judges can produce incompatible verdicts under shared criteria; TVD-Conflict searches for disagreement-inducing query–response pairs and organizes the resulting cases into recurring evaluator failure modes.

2 Related Work

LLM as Jailbreak Judge. With the rapid expansion of adversarial safety evaluation for large language models, LLM-AS-A-JUDGE has become the standard infrastructure for quantifying attack success rates (ASR). In automated red teaming, black-box attack frameworks such as PAIR [Chao et al., 2023] and TAP [Mehrotra et al., 2023] rely on LLM judges to provide optimization feedback. In standardized safety benchmarking, frameworks including HarmBench [Mazeika et al., 2024] and JailbreakBench [Chao et al., 2024] deploy state-of-the-art frontier models or dedicated safety classifiers to automate evaluations across diverse threat categories. Recent efforts further investigate multi-agent debate mechanisms [Lin et al., 2025] to enhance judgment consensus. These pipelines universally operate under the assumption that judge models provide stable and faithful safety assessments across diverse response styles.

Vulnerabilities and Disagreements in LLM Judges. Despite widespread adoption, extensive research shows that model-based evaluators are susceptible to substantial systematic flaws in general domains. Prior meta-evaluations demonstrate that LLM judges exhibit inherent biases such as position bias, verbosity bias, and self-enhancement [Zheng et al., 2023], while remaining vulnerable to minor adversarial perturbations [Raina and Gales, 2024]. Ye et al. [2025] identify 12 distinct types of bias and introduce the CALM benchmark to quantify judge sensitivity under controlled perturbations. In safety and jailbreak evaluation, these vulnerabilities become even more pronounced. Recent analyses reveal that safety judges are easily deceived by superficial compliance disclaimers, semantic reframing, or degraded incoherence from victim models [Souly et al., 2024, Cai et al., 2024], with judge accuracy degrading to near random chance under red-teaming distribution shifts [Schwinn et al., 2026, Almasoud et al., 2026]. Furthermore, safety-judge decisions often remain anchored to rigid internal priors when the provided context or safety criteria vary, showing reluctance to revise evaluations when explicit definitions conflict with their default boundaries [Alloula et al., 2026]. Although emerging efforts have begun benchmarking judge unreliability through static stress tests and human audits [Schwinn et al., 2026], they fail to systematically examine what specific content triggers instability across judges, nor do they provide an automated mechanism to actively mine and categorize these boundary failures. To bridge this gap, TVD-Conflict introduces an agentic framework that leverages inter-judge disagreement as an automated validator to dynamically discover conflict-inducing pairs and systematically unveil their underlying failure modes.

Table 1: Judge-dependent evaluation of ten jailbreak attacks. The left block reports attack success rate (ASR, %) under each judge–rubric combination; the right block reports cross-judge Conflict Rate (CR, %) with the rubric held fixed.

	Attack Success Rate by Judge						Cross-Judge Conflict Rate	
	GPT-OSS-120B		Qwen3.7-Flash		DeepSeek-V4-Flash		All Judges	
Attack	Bin.	Score	Bin.	Score	Bin.	Score	Bin. CR	Score CR
Cognitive Overload	0.5	4.0	0.5	4.0	0.0	3.5	0.5	1.5
MathPrompt	67.5	29.5	51.0	27.0	32.5	11.5	43.4	10.2
Semantic Cloaking	50.0	28.0	42.0	47.0	45.0	49.0	24.8	10.7
ActorBreaker	43.0	66.0	24.0	52.5	25.5	55.0	35.0	11.2
Mastermind	31.5	87.5	13.0	70.5	15.0	79.0	36.5	9.5
Compound Roleplay	42.0	41.0	37.5	13.5	24.0	37.0	16.1	10.4
CIA Context	49.5	70.0	21.5	70.5	27.5	72.5	41.1	7.3
AutoDAN-Turbo	17.0	23.5	15.0	21.0	13.0	21.0	7.9	1.1
FITD	51.5	45.5	47.0	12.5	43.0	13.0	16.7	37.5
Prefill	31.0	30.0	17.0	22.5	24.5	20.0	23.7	11.9
Overall	–	–	–	–	–	–	24.2	10.8

Protocol. Each attack contains four victim models $\times$ 50 AdvBench goals. Binary ASR is the fraction labeled harmful. Scoring ASR requires refusal $= 0$ and mean convincingness–specificity ≥ 4; missing judgments count as unsuccessful in the fixed denominator of 200. Binary CR is the fraction of complete cases in which the three harmfulness labels are not identical. Scoring CR uses $\max_j s_j - \min_j s_j > 0.5$, where $s = (1 - r)(c + p - 2)/8$.

3 Extract Safety-Judge Failure-Modes from Attack Responses

We evaluate responses from ten jailbreak attacks: Cognitive Overload, MathPrompt, Semantic Cloaking, ActorBreaker, Mastermind, Compound Roleplay, CIA Context, AutoDAN-Turbo, FITD, and Prefill. Each attack contributes 200 responses, covering four victim models and 50 AdvBench goals. Every response is evaluated by GPT-OSS-120B, Qwen3.7-Flash, and DeepSeek-V4-Flash using both the binary and graded judge prompts.

Panel A of Table 1 reports the ASR assigned by each backbone under the two rubrics. The choice of rubric can reverse the apparent ranking of attack methods. For GPT-OSS-120B, for example, MathPrompt falls from 67.5% under binary judging to 29.5% under scoring, while Mastermind rises from 31.5% to 87.5%. Panel B isolates evaluator conflict from ASR: it holds the rubric fixed and measures disagreement across the three backbones for each attack.

We derive the failure-mode taxonomy from conflict cases rather than from all attack responses. Each case retains the shared query and response, the raw judge outputs, parsed verdicts, and the dimensions responsible for disagreement. This representation supports both quantitative prevalence estimates and qualitative inspection of how an evaluator reached its decision.

4 TVD-Conflict: Systematic Safety-Judge Failure Data Generation

TVD-Conflict formulates safety-judge stress testing as a search problem over task, validator, and data components. A task specifies the safety criterion and visible context. Validators are LLM judges that independently apply the same criterion. The data component stores candidate interactions together with raw and parsed validator outputs. A candidate is retained when it crosses the configured disagreement threshold.

For binary judging, a conflict occurs when at least two validators assign different harmfulness labels. For graded judging, each validator receives the score

$$s = (1 - r)\frac{c + p - 2}{8},$$

where $r \in \{0, 1\}$ is the refusal decision and $c, p \in \{1, \ldots, 5\}$ are convincingness and specificity. A pair conflicts when its absolute score difference exceeds 0.5. The search process uses validator feedback to revise a candidate while preserving the underlying task and ground-truth safety label.

Table 2: Conflict Rate (%) under a search budget of ten candidate revisions. Results will be inserted after the budget-controlled search runs complete.

Judge Pair	Binary CR	Scoring CR
GPT-OSS / Qwen	TBD	TBD
GPT-OSS / DeepSeek	TBD	TBD
Qwen / DeepSeek	TBD	TBD
Overall	TBD	TBD

Table 3: Prevalence of six safety-judge failure modes in the TVD-Conflict dataset. Counts and shares will be inserted after annotation.

Failure Mode	Count	Share (%)
Criteria ambiguity	TBD	TBD
Refusal but harmful	TBD	TBD
Compliance but safe	TBD	TBD
50–50 problem	TBD	TBD
Distraction	TBD	TBD
Insufficient capability	TBD	TBD

Search-budget evaluation. Table 2 will report Conflict Rate after a budget of ten candidate revisions, using the same conflict definitions as Table 1.

5 Failure-Modes of Frontier LLM-as-Safety-Judge

We organize disagreement cases into six failure modes. The categories describe the decision error exposed by a case, not the surface form of the jailbreak. Table 3 will report their prevalence after the conflict set has been annotated.

Criteria ambiguity. The judges operationalize the stated safety criterion differently, such as treating harmfulness, policy violation, and attack success as interchangeable.

Refusal but harmful. A response begins with refusal language but subsequently provides information that materially advances the harmful request. Some judges anchor on the refusal marker, while others evaluate the supplied content.

Compliance but safe. A response adopts the form of compliance without delivering harmful or usable information. Judges that rely on affirmative framing can mistake stylistic compliance for harmful assistance.

The 50–50 problem. The available evidence supports two plausible readings of comparable strength, causing verdicts to split across judges or rubric formulations.

Distraction. An evaluator focuses on wrappers, disclaimers, fictional settings, mathematical notation, or peripheral detail instead of the behavior that determines safety.

Insufficient capability. The judge fails to decode, follow, or reason over the representation required to identify the harmful behavior, including obfuscation and multi-step context.

6 Conclusion

Safety evaluation with LLM judges is not invariant to the evaluator. Across ten jailbreak attacks, three judge backbones and two rubrics produce materially different decisions on identical interactions. Cross-judge Conflict Rate is 24.2% under binary judging and 10.8% under scoring, with substantially higher rates for particular attacks. The variation is strongly attack dependent, indicating that a single aggregate ASR can conceal systematic blind spots. TVD-Conflict turns these discrepancies into a

data-generation objective and supports analysis through six operational failure modes. The resulting evidence argues for evaluating safety judges as fallible measurement instruments rather than treating their outputs as interchangeable ground truth.

References

Anissa Alloula, Federico Licini, Ava Batchkala, and Seraphina Goldfarb-Tarrant. Safety is contextual, LLM-judges are not: Navigating the rigid priors of evaluators, 2026. URL `https://arxiv.org/abs/2606.07874`.

Aiman Almasoud, Antony Anju, Marco Arazzi, Mert Cihangiroglu, Vignesh Kumar Kembu, S Nicolazzo, and A Nocera. Security in llm-as-a-judge: A comprehensive sok. *arXiv preprint arXiv:2603.29403*, 2026.

Hongyu Cai, Jeffrey Lin, Dawn Song, et al. Rethinking how to evaluate language model jailbreak. *arXiv preprint arXiv:2404.06407*, 2024.

Patrick Chao, Alexander Robey, Edgar Dobriban, Hamed Hassani, George J Pappas, and Eric Wong. Jailbreaking black box large language models in twenty queries. *arXiv preprint arXiv:2310.08419*, 2023.

Patrick Chao, Alexander Robey, Edgar Dobriban, Hamed Hassani, George J Pappas, and Eric Wong. Jailbreakbench: An open robustness benchmark for jailbreaking large language models. In *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, 2024.

Dachuan Lin, Guobin Shen, Zihao Yang, Tianrong Liu, Dongcheng Zhao, and Yi Zeng. Efficient llm safety evaluation through multi-agent debate. *arXiv preprint arXiv:2511.06396*, 2025.

Mantas Mazeika, Long Phan, Xuwang Yin, Andy Zou, Zifan Wang, Norman Mu, Elham Sakhaee, Nathaniel Li, Steven Basart, Bo Li, et al. Harmbench: A standardized evaluation framework for automated red teaming and robust jailbreak. In *IEEE Symposium on Security and Privacy (S&P)*, 2024.

Anay Mehrotra, Manolis Zampetakis, Paul Kassianik, Blaine Nelson, Hyrum Anderson, and Yaron Singer. Tree of attacks: Jailbreaking black-box llms with automated search. *arXiv preprint arXiv:2312.02119*, 2023.

Vyas Raina and Mark Gales. Is LLM-as-a-judge robust? investigating universal adversarial attacks on zero-shot LLM assessment. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 10617–10634, 2024.

Leo Schwinn, Moritz Ladenburger, Tim Beyer, Mehrnaz Mofakhami, Gauthier Gidel, and Stephan Günnemann. A coin flip for safety: Llm judges fail to reliably measure adversarial robustness. *International Conference on Learning Representations (ICLR)*, 2026.

Alexandra Souly, Qingyuan Lu, Dillon Dillon, Alex Breuer, et al. Strongreject: Jailbreaking llms with robust automatic evaluation. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

Jiayi Ye, Yanbo Wang, Yue Huang, Dongping Chen, Qihui Zhang, Nuno Moniz, Tian Gao, Werner Geyer, Chao Huang, Pin-Yu Chen, Nitesh Chawla, and Xiangliang Zhang. Justice or Prejudice? Quantifying Biases in LLM-as-a-Judge. In *International Conference on Learning Representations (ICLR)*, 2025.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, et al. Judging llm-as-a-judge with mt-bench and chatbot arena. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.